


Deep Learning

TP2 : Generative models

SIA_TP1
TRAN Jeremy

1 : Autoencoder




The process is to compress the input then reconstruct it to learn its features :

- The encoder compresses the data into a latent space, here using down-samplers, adding more of these layers creates a smaller latent space.

- The decoder reconstructs the image from the latent space data using up-samplers.

$$\text{loss} = ||x - \hat{x}||^2 = ||x - d(z)||^2 = ||x - d(e(x))||^2$$

train_loss



Run ↑

Smoothed

Value

Step

Relative

● lightning_logs\version_35


0.0151

0.0123

4,499

1.389 min

val_loss



Run ↑

Smoothed

Value

Step

Relative ↑

● lightning_logs\version_35

0.0141

0.0083

4,499

59.77 sec


Some results :
Depth 3 AE, 3 epochs

Test metric

DataLoader 0

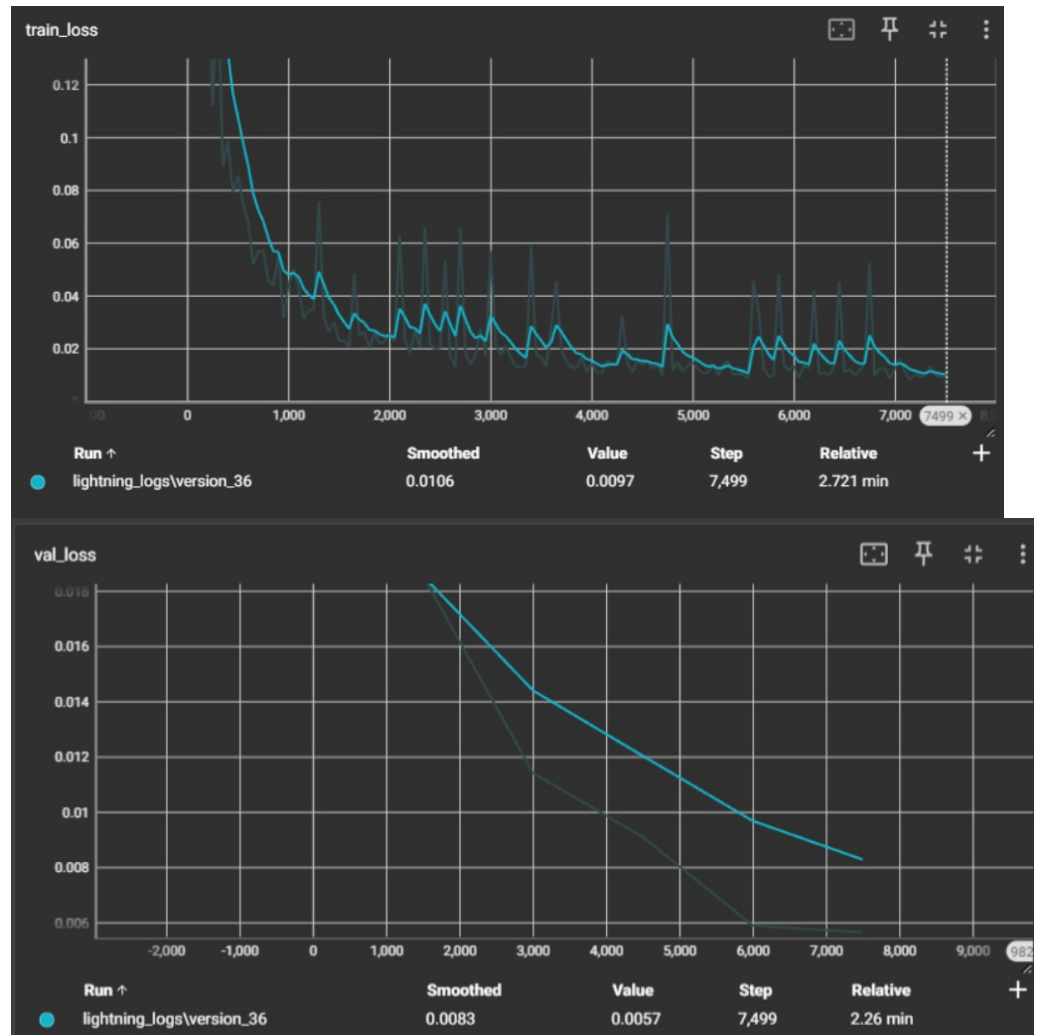
test_loss

0.00817732885479927



The resulting
images are very
close to the originals

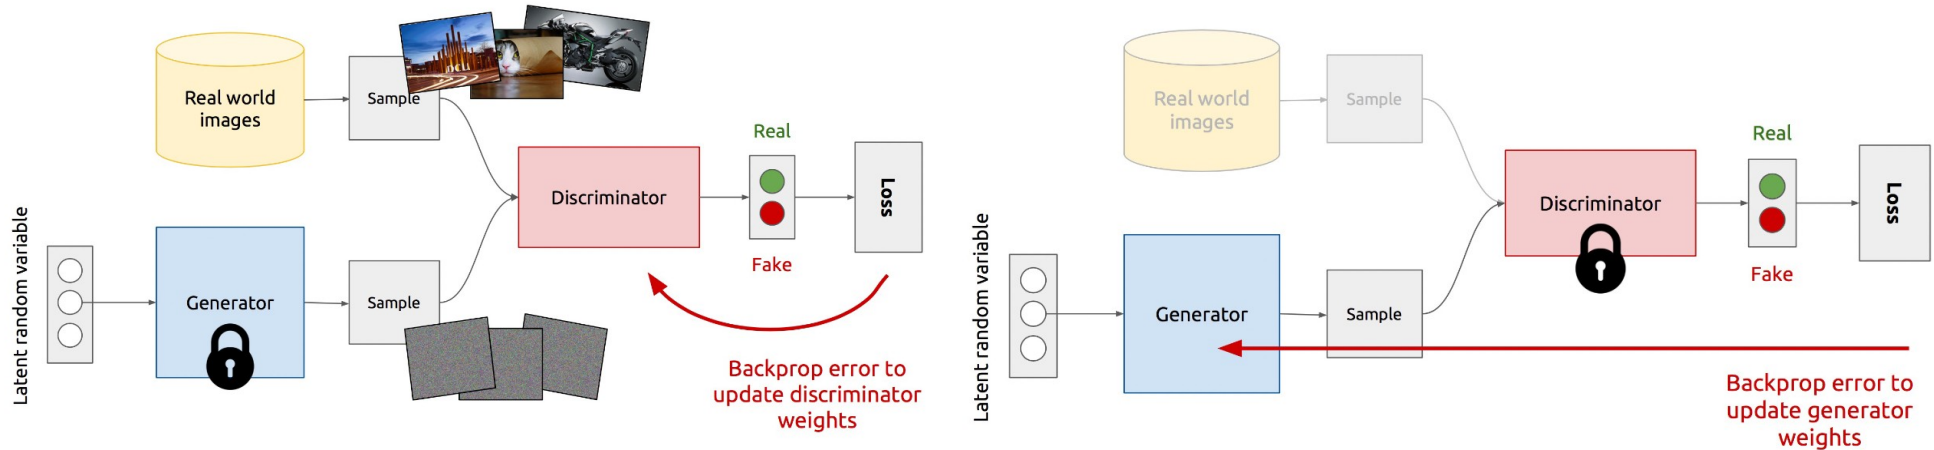
Test using different training parameters :



Here the depth is 5 and trained on 5 epochs.

There are visible defects in the image not present in the previous test.

2 : Generative adversarial network



The GAN model uses a generator to create fake images and a discriminator to distinguish between fake and real images.

- In the first step, the discriminator is trained to recognize real pictures
- In the second step, the generator that used to generate random images has its weights updated to output images that gradually get closer to the real images.

```

model = GAN()
model

GAN(
  (generator): Generator(
    (model): AutoEncoder2(
      (encoder): Encoder(
        (encoder): ModuleList(
          (0): ConvDown(
            (model): Sequential(
              (0): Conv2d(3, 8, kernel_size=(3, 3), stride=(1, 1))
              (1): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
              (2): Dropout2d(p=0.5, inplace=False)
              (3): LeakyReLU(negative_slope=0.2)
            )
          )
          (1): ConvDown(
            (model): Sequential(
              (0): Conv2d(8, 16, kernel_size=(3, 3), stride=(1, 1))
              (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
              (2): Dropout2d(p=0.5, inplace=False)
              (3): LeakyReLU(negative_slope=0.2)
            )
          )
          (2): ConvDown(
            (model): Sequential(
              (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
              (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            )
          )
          ...
          (3): Conv2d(128, 1, kernel_size=(3, 3), stride=(1, 1), bias=False)
        )
      )
    )
  )
)

```

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings...](#)

I managed to create the GAN model but a bug probably due to some update prevented me from actually testing it.